Authoritative Three-File Corpus Audit

Scope: only the three permitted current files were used. No legacy datasets were read or merged. Jung Chang is excluded from L1.

Audit item	Observed
Workbook L1 names	41; 40 after Jung Chang exclusion
L1 author matches in corpus CSV	40/40
Eligible L1 texts (>=100 tokens)	515 texts, 37 authors
Workbook L2 names	45
L2 sample corpus	2 essays; no author IDs
Inferential tests	BLOCKED: both groups must have >=40 identifiable writers

L1 records include abstracts and empty-text records; a text record is not necessarily comparable authored prose. The two anonymous L2 essays cannot stand in for 45 named writers. Accordingly, this report contains coverage/audit and descriptive metrics only. No inferential p-values or effect sizes are claimed.

Metrics implemented for eligible texts: token count, mean word length, type-token ratio (TTR), and Carroll CTTR. Author-level summaries aggregate eligible texts. See README for thresholds and full limitations.